

Unsupervised Learning: Clustering

Valeria Vitelli

Oslo Centre for Biostatistics and Epidemiology

Department of Biostatistics, UiO

valeria.vitelli@medisin.uio.no

MED3007

Statistical Principles in Genomics: an Introduction with Rstudio

16.01.2025

- 1 Introduction
 - Motivation
 - Notions of similarity
- 2 Hierarchical Clustering Algorithms
 - Main ideas
 - Illustrative example
 - Visualization
 - Summary
- 3 Partitioning clustering: k -means
 - Partitioning Clustering Algorithms
 - Main idea of the k -means
 - Illustrative example 1
 - Illustrative example 2
- 4 Summary

Recap from last time...

Unsupervised methods


Methods that **do not make use of external information**: groups / class assignments, clinical outcomes, covariates.

Aim: finding hidden structure in the data

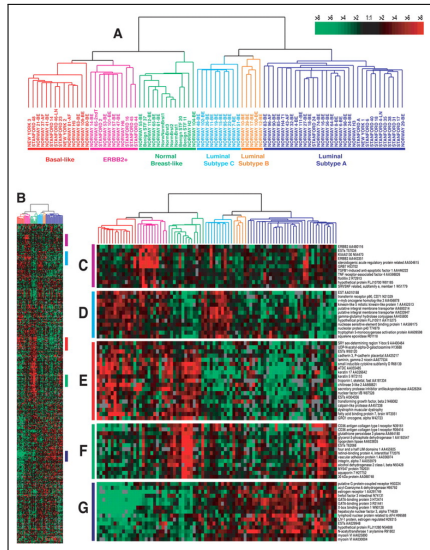
How can we find and/or visualise structure?

- **Reducing dimensionality** by removing noise and “uninteresting” stuff
→ TODAY – Lecture 2
- **Ordering and grouping** together via clustering
→ TOMORROW – Lecture 3

Heatmaps are everywhere...



Eisen et al. (1998), Figure 1



Sorlie et al. (2001), Figure 1

What is clustering?

Clustering methods

Methods that **aim at grouping** a collection of objects into **groups**, or clusters, such that objects **within each cluster** are more closely related to one another than objects assigned to a **different cluster**

- **Distance measure**

A notion of distance or similarity of two objects: When are two objects close to each other?

- **Clustering algorithm**


A procedure to minimise distances of objects within groups and/or maximise distances between groups.

What is meant when two patients are said to be “similar”?

Possibilities:

- Red and black patients are similar:
They lie close to each other.
- Blue and black patients are similar:
They are positively correlated.
- Red and blue patients are associated:
They are negatively correlated.

Expressions of 3 patients at 4 genes



Examples of distance measures $d(\cdot, \cdot)$

- Euclidean distance

$$d_{\text{euc}}(\mathbf{x}, \mathbf{y}) = \left(\sum_{k=1}^m (x_k - y_k)^2 \right)^{1/2}$$

- 1 - Pearson's correlation

$$d_{\text{cor}}(\mathbf{x}, \mathbf{y}) = 1 - r(\mathbf{x}, \mathbf{y}) = 1 - \frac{\sum_{i=1}^m (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^m (x_i - \bar{x})^2 \sum_{i=1}^m (y_i - \bar{y})^2}}$$

- 1 - Spearman's rank correlation ($R(x_i)$ = Rank of x_i)

$$\begin{aligned} d_{\text{spear}}(\mathbf{x}, \mathbf{y}) &= 1 - r_s(\mathbf{x}, \mathbf{y}) \\ &= 1 - \frac{\sum_{i=1}^m (R(x_i) - \overline{R(x)})(R(y_i) - \overline{R(y)})}{\sqrt{\sum_{i=1}^m (R(x_i) - \overline{R(x)})^2 \sum_{i=1}^m (R(y_i) - \overline{R(y)})^2}} \end{aligned}$$

Examples of distance measures $d(\cdot, \cdot)$ – visual!

Euclidean distance:


- $d(\text{black}, \text{red}) = 4.90$
- $d(\text{black}, \text{blue}) = 6.00$
- $d(\text{blue}, \text{red}) = 9.16$

Spearman-correlation-distance:

- $d(\text{black}, \text{red}) = 2.00$
- $d(\text{black}, \text{blue}) = 0.00$
- $d(\text{blue}, \text{red}) = 2.00$


Interpretation:

- the smaller the distance, the more similar the patients' response
- different distances measure “similarity” differently



- 1 Introduction
 - Motivation
 - Notions of similarity
- 2 Hierarchical Clustering Algorithms
 - Main ideas
 - Illustrative example
 - Visualization
 - Summary
- 3 Partitioning clustering: k -means
 - Partitioning Clustering Algorithms
 - Main idea of the k -means
 - Illustrative example 1
 - Illustrative example 2
- 4 Summary

Hierarchical clustering - Agglomerative algorithm



- **Bottom-up algorithm** (top-down methods are less common)
- Start with each object assigned to its own cluster.
- In each iteration, merge the two clusters with the minimal distance from each other - until you are left with a single cluster containing all objects
- But how define the distance **between two clusters**?

Hierarchical clustering - Linkage

Calculation of distance $d(G, H)$ between clusters G and H is based on pairwise distances between objects from the two clusters:

- **Single linkage** uses the smallest distance between the objects:

$$d_S(G, H) = \min_{(i \in G; j \in H)} d_{ij}$$

Single linkage is not commonly used.

- **Complete linkage** uses the largest distance between the objects:


$$d_C(G, H) = \max_{(i \in G; j \in H)} d_{ij}$$

- **Average linkage** uses the average distance between the objects:


$$d_A(G, H) = \frac{1}{N_G N_H} \sum_{i \in G} \sum_{j \in H} d_{ij}$$

Hierarchical clustering - Linkage illustration


Single linkage



Complete linkage



Average linkage



Hierarchical clustering - Single linkage example

- Data: 5 patients 1-5 (= rows), 2 expression arrays A,B (= columns)
- Method: **Agglomerative hierarchical clustering** using **Single-Linkage**


Expression matrix

	A	B
1	5.0	7.0
2	5.5	8.5
3	8.0	8.5
4	8.0	7.0
5	8.0	5.0


Distance matrix
(**Euclidean distance**)

	1	2	3	4	5
1	0.00	2.50	11.25	9.00	13.00
2	2.50	0.00	6.25	8.50	18.50
3	11.25	6.25	0.00	2.25	12.25
4	9.00	8.50	2.25	0.00	4.00
5	13.00	18.50	12.25	4.00	0.00


Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example



Agglomerative clustering - Single linkage example




Agglomerative clustering - Single linkage example




Hierarchical clustering - Linkage methods


Same distance matrix!



Single Linkage




Complete Linkage



Average Linkage

Hierarchical clustering - Heatmap



- Single linkage
- Euclidean distance
- The expression values are represented as colours.

Summary & Take-home messages

- The procedure provides a **hierarchy of the clustering**, with the number of clusters ranging from 1 to the number of objects
- **Finding a meaningful cut** is similar to the problem of finding the number of clusters k for partitioning methods (NEXT SLIDES)
- An incorrect merge early in the tree **cannot be changed later on**
- The **choice of the distance measure** depends on the data and the intention of the clustering
- Even data generated at random will result in a clustering: be careful with interpretation!

- 1 Introduction
 - Motivation
 - Notions of similarity
- 2 Hierarchical Clustering Algorithms
 - Main ideas
 - Illustrative example
 - Visualization
 - Summary
- 3 Partitioning clustering: k -means
 - Partitioning Clustering Algorithms
 - Main idea of the k -means
 - Illustrative example 1
 - Illustrative example 2
- 4 Summary

Partitioning

Partitioning algorithms (= non-hierarchical methods)

- split the data into a **pre-specified number k** of groups
 - **iteratively re-allocate** objects until some **optimality criterion** is met
- k thus needs to be **fixed in advance**

Examples:

- **k -means clustering**
- **Partitioning around medoids (PAM)**
generalization of k -means (allows additional optimisation criteria)
- **Self-organising maps (SOM)**
similar to k -means but with additional constraints (grid-like structure)

Partitioning - k -means

k -means


The k -means algorithm minimises the sum of within-cluster variances

It chooses a random sample of k different objects as **initial cluster centroids**, then alternating until convergence:

- 1 Assign each object to the cluster whose centroid is closest (among the k centroids) with respect to Euclidean distance
- 2 Calculate k new centroids as the averages of all points assigned to the same cluster

Partitioning - k -means example

Start




$$k = 3$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - *k*-means example




Centroid matrix

$$\begin{pmatrix} 5.0 & 6.0 \\ 7.0 & 7.5 \\ 6.0 & 5.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - k -means example



Centroid matrix

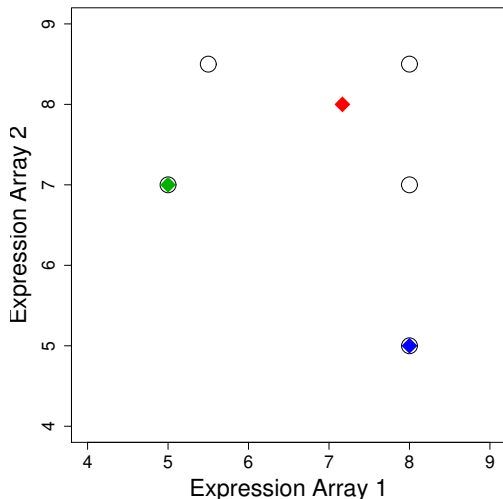
$$\begin{pmatrix} 5.0 & 6.0 \\ 7.0 & 7.5 \\ 6.0 & 5.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - *k*-means example

Iteration 1



Centroid matrix


$$\begin{pmatrix} 5.0 & 7.0 \\ 7.17 & 8.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - *k*-means example

Iteration 1




Centroid matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 7.17 & 8.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - *k*-means example

Iteration 2

Centroid matrix


$$\begin{pmatrix} 5.25 & 7.75 \\ 8.00 & 7.75 \\ 8.0 & 5.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - *k*-means example

Iteration 2




Centroid matrix

$$\begin{pmatrix} 5.25 & 7.75 \\ 8.00 & 7.75 \\ 8.0 & 5.0 \end{pmatrix}$$

Expression matrix


$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - k -means example



Iteration 3: No changes in centroid matrix

Partitioning - k -means example



Iteration 3: No changes
→ We are done.


Partitioning - k -means

Choice of cluster number and initial starting values for cluster centroids:

- Changing the **number of clusters** can completely change the cluster structure. This is contrary to hierarchical clustering.
- k -means is a **randomised algorithm**: two runs usually produce different results. Thus, it has to be applied several times and the result with minimal sum of within-cluster-variances should be chosen.
Even when doing so, we are not guaranteed to find the **best solution**.

Example: same data set (5 patients, 2 expression values)

Method: **Partitioning clustering** with k -means with $k = 2$

Partitioning - *k*-means example 2

Centroid matrix


$$\begin{pmatrix} 4.5 & 8.0 \\ 7.0 & 5.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - k -means example 2

Iteration 1




Centroid matrix

$$\begin{pmatrix} 6.17 & 8.0 \\ 8.0 & 6.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - *k*-means example 3

Centroid matrix


$$\begin{pmatrix} 8.0 & 4.0 \\ 7.0 & 9.0 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

Partitioning - k -means example 3

Iteration 1



Centroid matrix

$$\begin{pmatrix} 8.0 & 5.0 \\ 6.63 & 7.75 \end{pmatrix}$$

Expression matrix

$$\begin{pmatrix} 5.0 & 7.0 \\ 5.5 & 8.5 \\ 8.0 & 8.5 \\ 8.0 & 7.0 \\ 8.0 & 5.0 \end{pmatrix}$$

- 1 Introduction
 - Motivation
 - Notions of similarity
- 2 Hierarchical Clustering Algorithms
 - Main ideas
 - Illustrative example
 - Visualization
 - Summary
- 3 Partitioning clustering: k -means
 - Partitioning Clustering Algorithms
 - Main idea of the k -means
 - Illustrative example 1
 - Illustrative example 2
- 4 Summary

Clustering - Summary

Hierarchical clustering:

- Procedure provides a **hierarchy of the clustering**, with the number of clusters ranging from 1 to the number of objects
- Incorrect merges early in the tree **cannot be changed later on**
- Careful with interpretation of dendrograms and resulting heatmaps!
The order of objects within a cluster is arbitrary

Partitioning algorithms:

- Careful with initialization...
- How to choose the correct number of groups?

Clustering - Summary

All clustering methods:

- The **choice of the distance measure** depends on the data and the intention of the clustering
- Use objective measures to support the decision for number of clusters
- Even data **generated at random** will result in clusters
- Be careful with pre-selection of features before clustering!

Practical Issues in Clustering

- Some **decisions** to be made:
 - Should the variables be standardized first? Other transformations needed to achieve normally distributed data?
 - Do we need a pre-selection of variables? This needs to be unsupervised, i.e. using variances not t-tests
- In case of **hierarchical clustering**:
 - What dissimilarity measure should be used?
 - What type of linkage should be used?
 - Where should we cut the dendrogram in order to obtain clusters?
- In case of **partitioning clustering**:
 - How many clusters should we look for?
- **Try several choices**, and look for clustering results that are most useful for interpretation. There is no single right answer!

R/ Bioconductor

Heatmaps

- `heatmap()` (stats) and many enhancements, e.g. `pheatmap()` (pheatmap) and `Heatmap()` (ComplexHeatmap)

Hierarchical clustering

- `hclust()` (stats), `hcluster` (amap)

Partitioning clustering

- `kmeans()` (stats), `pam()` (cluster)

CRAN Task View Cluster

<https://cran.r-project.org/web/views/Cluster.html>